



Contents

Fine-Tuning for LLMs: from Beginner to Advanced

Challenge: Fine-tuning the Q&A model

Leave a review



397



9,765



Challenge: Fine-Tuning the Q&A model

Introduction:

Ready for an exciting challenge? This time, you'll fine-tune a Question & Answer (Q&A) model using the FLAN-T5-large model and the SQuAD dataset, leveraging transfer learning. This exercise will help you adapt a powerful pre-trained model to excel at answering questions based on given passages, an essential skill for building advanced NLP solutions.

Steps:

1. Load data:

- Download and preprocess the SQuAD dataset to prepare it for training.

2. Initialize model:

- Load the pre-trained FLAN-T5-large model and tokenizer from Hugging Face's Transformers library.

3. Do transfer learning:

- Freeze the model's layers to enable efficient fine-tuning with limited data.

4. Prepare data for training:

- Tokenize the dataset and create training and validation splits.

5. Fine-tune the model:

- Train the FLAN-T5-large model on the tokenized dataset, adjusting only the parameters of the LoRA adapters.

6. Evaluate performance:

- Assess the model's performance using metrics such as BLUE or ROUGE score, Exact Match to ensure it accurately answers questions.

Conclusion:

By completing this challenge, you've learned to fine-tune a Q&A model using FLAN-T5-large, an essential component of a comprehensive NLP solution. This ties together with sentiment analysis and translation to create a robust system capable of understanding, interpreting, and responding to human language in diverse applications.

[Previous](#)[Next](#)